

Animal Image Classifier using CNN

1. Data Preparation & Exploration

- Download and extract the dataset (from Kaggle or your local folder).
- Organize images into train/validation/test folders by class.
- Visualize sample images from each class.
- Perform data augmentation to increase robustness.. Data Preparation & Exploration

```
import os
import shutil
import random
from glob import glob
from tensorflow.keras.preprocessing.image import ImageDataGenerator,
load_img, img_to_array

import matplotlib.pyplot as plt

# Set paths
base_dir = 'dataset/Animals'
classes = ['cats', 'dogs', 'snakes']
output_base = 'dataset/organized'

# Create train/val/test split
split_ratios = {'train': 0.8, 'val': 0.1, 'test': 0.1}

# Organize images into train/val/test folders
if not os.path.exists(output_base):
    for cls in classes:
        img_paths = glob(os.path.join(base_dir, cls, '*'))
        random.shuffle(img_paths)
        n_total = len(img_paths)
        n_train = int(n_total * split_ratios['train'])
        n_val = int(n_total * split_ratios['val'])
        splits = {
            'train': img_paths[:n_train],
            'val': img_paths[n_train:n_train + n_val],
            'test': img_paths[n_train + n_val:]
        }
        for split, files in splits.items():
            split_dir = os.path.join(output_base, split, cls)
            os.makedirs(split_dir, exist_ok=True)
            for f in files:
                shutil.copy(f, split_dir)

# Visualize sample images from each class
fig, axes = plt.subplots(1, len(classes), figsize=(12, 4))
for idx, cls in enumerate(classes):
    sample_dir = os.path.join(output_base, 'train', cls)
```

```

    sample_img = random.choice(os.listdir(sample_dir))
    img_path = os.path.join(sample_dir, sample_img)
    img = load_img(img_path, target_size=(128, 128))
    axes[idx].imshow(img)
    axes[idx].set_title(cls)
    axes[idx].axis('off')
plt.tight_layout()
plt.show()

# Data augmentation
datagen = ImageDataGenerator(
    rotation_range=20,
    width_shift_range=0.1,
    height_shift_range=0.1,
    shear_range=0.1,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)

# Visualize augmented images for one sample
sample_dir = os.path.join(output_base, 'train', classes[0])
sample_img = random.choice(os.listdir(sample_dir))
img_path = os.path.join(sample_dir, sample_img)
img = load_img(img_path, target_size=(128, 128))
x = img_to_array(img)
x = x.reshape((1,) + x.shape)

fig, ax = plt.subplots(1, 5, figsize=(15, 3))
i = 0
for batch in datagen.flow(x, batch_size=1):
    ax[i].imshow(batch[0].astype('uint8'))
    ax[i].axis('off')
    i += 1
    if i == 5:
        break
plt.suptitle(f'Augmented samples: {classes[0]}')
plt.show()

```



Augmented samples: cats



2. Data Loading with Keras

- Use `ImageDataGenerator` for loading and augmenting images.
- Split data into training, validation, and test sets.

```
# Prepare ImageDataGenerators for train, validation, and test sets
train_dir = os.path.join(output_base, 'train')
val_dir = os.path.join(output_base, 'val')
test_dir = os.path.join(output_base, 'test')

img_height, img_width = 128, 128 # Standard for transfer learning
batch_size = 32 # batch_size means number of images processed
together

# Training data generator with augmentation
train_gen = datagen.flow_from_directory(
    train_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
    class_mode='categorical'
)

# Validation and test generators without augmentation
test_val_datagen = ImageDataGenerator()
val_gen = test_val_datagen.flow_from_directory(
    val_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
```

```

        class_mode='categorical',
        shuffle=False
    )
    test_gen = test_val_datagen.flow_from_directory(
        test_dir,
        target_size=(img_height, img_width),
        batch_size=batch_size,
        class_mode='categorical',
        shuffle=False
    )

```

```

Found 2100 images belonging to 3 classes.
Found 450 images belonging to 3 classes.
Found 450 images belonging to 3 classes.

```

3. Load Pre-trained ResNet Model

- Use ResNet50 (or ResNet101 for more complexity) with imagenet weights.
- Exclude the top layer (include_top=False) to add custom layers.

```

from tensorflow.keras.applications import ResNet50

# Load ResNet50 with imagenet weights, excluding the top layer
base_model = ResNet50(weights='imagenet', include_top=False,
input_shape=(128, 128, 3))
base_model.trainable = False # Freeze base model layers

base_model.summary()

```

Model: "resnet50"

Layer (type)	Output Shape	Param #	Connected to
input_layer_1 (InputLayer)	(None, 128, 128, 3)	0	-
conv1_pad input_layer_1[0]... (ZeroPadding2D)	(None, 134, 134, 3)	0	
conv1_conv (Conv2D)	(None, 64, 64, 3)	9,472	conv1_pad[0]

[0]		64)		
	conv1_bn	(None, 64, 64,	256	conv1_conv[0]
[0]	(BatchNormalizatio...	64)		
	conv1_relu	(None, 64, 64,	0	conv1_bn[0]
[0]	(Activation)	64)		
	pool1_pad	(None, 66, 66,	0	conv1_relu[0]
[0]	(ZeroPadding2D)	64)		
	pool1_pool	(None, 32, 32,	0	pool1_pad[0]
[0]	(MaxPooling2D)	64)		
	conv2_block1_1_conv	(None, 32, 32,	4,160	pool1_pool[0]
[0]	(Conv2D)	64)		
	conv2_block1_1_bn	(None, 32, 32,	256	
conv2_block1_1_c...	(BatchNormalizatio...	64)		
	conv2_block1_1_relu	(None, 32, 32,	0	
conv2_block1_1_b...	(Activation)	64)		
	conv2_block1_2_conv	(None, 32, 32,	36,928	
conv2_block1_1 r...				

(Conv2D)	64)		
conv2_block1_2_bn	(None, 32, 32,	256	
conv2_block1_2_c...	(BatchNormalizatio...	64)	
conv2_block1_2_relu	(None, 32, 32,	0	
conv2_block1_2_b...	(Activation)	64)	
conv2_block1_0_conv	(None, 32, 32,	16,640	pool1_pool[0]
[0]	(Conv2D)	256)	
conv2_block1_3_conv	(None, 32, 32,	16,640	
conv2_block1_2_r...	(Conv2D)	256)	
conv2_block1_0_bn	(None, 32, 32,	1,024	
conv2_block1_0_c...	(BatchNormalizatio...	256)	
conv2_block1_3_bn	(None, 32, 32,	1,024	
conv2_block1_3_c...	(BatchNormalizatio...	256)	
conv2_block1_add	(None, 32, 32,	0	
conv2_block1_0_b...	(Add)	256)	
conv2_block1_3_b...			
conv2_block1_out	(None, 32, 32,	0	
conv2_block1_add...	(Activation)	256)	

conv2_block2_1_conv conv2_block1_out...	(None, 32, 32, (Conv2D)	16,448 64)	
conv2_block2_1_bn conv2_block2_1_c...	(None, 32, 32, (BatchNormalizatio...	256 64)	
conv2_block2_1_relu conv2_block2_1_b...	(None, 32, 32, (Activation)	0 64)	
conv2_block2_2_conv conv2_block2_1_r...	(None, 32, 32, (Conv2D)	36,928 64)	
conv2_block2_2_bn conv2_block2_2_c...	(None, 32, 32, (BatchNormalizatio...	256 64)	
conv2_block2_2_relu conv2_block2_2_b...	(None, 32, 32, (Activation)	0 64)	
conv2_block2_3_conv conv2_block2_2_r...	(None, 32, 32, (Conv2D)	16,640 256)	
conv2_block2_3_bn conv2_block2_3_c...	(None, 32, 32, (BatchNormalizatio...	1,024 256)	

conv2_block2_add	(None, 32, 32,	0	
conv2_block1_out...			
(Add)	256)		
conv2_block2_3_b...			
conv2_block2_out	(None, 32, 32,	0	
conv2_block2_add...			
(Activation)	256)		
conv2_block3_1_conv	(None, 32, 32,	16,448	
conv2_block2_out...			
(Conv2D)	64)		
conv2_block3_1_bn	(None, 32, 32,	256	
conv2_block3_1_c...			
(BatchNormalizatio...	64)		
conv2_block3_1_relu	(None, 32, 32,	0	
conv2_block3_1_b...			
(Activation)	64)		
conv2_block3_2_conv	(None, 32, 32,	36,928	
conv2_block3_1_r...			
(Conv2D)	64)		
conv2_block3_2_bn	(None, 32, 32,	256	
conv2_block3_2_c...			
(BatchNormalizatio...	64)		
conv2_block3_2_relu	(None, 32, 32,	0	
conv2_block3_2_b...			
(Activation)	64)		

conv2_block3_3_conv	(None, 32, 32,	16,640
conv2_block3_2_r...	(Conv2D)	256)
conv2_block3_3_bn	(None, 32, 32,	1,024
conv2_block3_3_c...	(BatchNormalizatio...	256)
conv2_block3_add	(None, 32, 32,	0
conv2_block2_out...	(Add)	256)
conv2_block3_3_b...		
conv2_block3_out	(None, 32, 32,	0
conv2_block3_add...	(Activation)	256)
conv3_block1_1_conv	(None, 16, 16,	32,896
conv2_block3_out...	(Conv2D)	128)
conv3_block1_1_bn	(None, 16, 16,	512
conv3_block1_1_c...	(BatchNormalizatio...	128)
conv3_block1_1_relu	(None, 16, 16,	0
conv3_block1_1_b...	(Activation)	128)
conv3_block1_2_conv	(None, 16, 16,	147,584
conv3_block1_1_r...	(Conv2D)	128)

conv3_block1_2_bn	(None, 16, 16,	512	
conv3_block1_2_c...	(BatchNormalizatio...	128)	
conv3_block1_2_relu	(None, 16, 16,	0	
conv3_block1_2_b...	(Activation)	128)	
conv3_block1_0_conv	(None, 16, 16,	131,584	
conv2_block3_out...	(Conv2D)	512)	
conv3_block1_3_conv	(None, 16, 16,	66,048	
conv3_block1_2_r...	(Conv2D)	512)	
conv3_block1_0_bn	(None, 16, 16,	2,048	
conv3_block1_0_c...	(BatchNormalizatio...	512)	
conv3_block1_3_bn	(None, 16, 16,	2,048	
conv3_block1_3_c...	(BatchNormalizatio...	512)	
conv3_block1_add	(None, 16, 16,	0	
conv3_block1_0_b...	(Add)	512)	
conv3_block1_3_b...			
conv3_block1_out	(None, 16, 16,	0	
conv3_block1_add...	(Activation)	512)	
conv3_block2_1_conv	(None, 16, 16,	65,664	

conv3_block1_out...	(Conv2D)	128)		
conv3_block2_1_bn	(None, 16, 16,	512		
conv3_block2_1_c...	(BatchNormalizatio...	128)		
conv3_block2_1_relu	(None, 16, 16,	0		
conv3_block2_1_b...	(Activation)	128)		
conv3_block2_2_conv	(None, 16, 16,	147,584		
conv3_block2_1_r...	(Conv2D)	128)		
conv3_block2_2_bn	(None, 16, 16,	512		
conv3_block2_2_c...	(BatchNormalizatio...	128)		
conv3_block2_2_relu	(None, 16, 16,	0		
conv3_block2_2_b...	(Activation)	128)		
conv3_block2_3_conv	(None, 16, 16,	66,048		
conv3_block2_2_r...	(Conv2D)	512)		
conv3_block2_3_bn	(None, 16, 16,	2,048		
conv3_block2_3_c...	(BatchNormalizatio...	512)		
conv3_block2_add	(None, 16, 16,	0		
conv3_block1_out...				

(Add)	512)		
conv3_block2_3_b...			
conv3_block2_out	(None, 16, 16,	0	
conv3_block2_add...			
(Activation)	512)		
conv3_block3_1_conv	(None, 16, 16,	65,664	
conv3_block2_out...			
(Conv2D)	128)		
conv3_block3_1_bn	(None, 16, 16,	512	
conv3_block3_1_c...			
(BatchNormalizatio...	128)		
conv3_block3_1_relu	(None, 16, 16,	0	
conv3_block3_1_b...			
(Activation)	128)		
conv3_block3_2_conv	(None, 16, 16,	147,584	
conv3_block3_1_r...			
(Conv2D)	128)		
conv3_block3_2_bn	(None, 16, 16,	512	
conv3_block3_2_c...			
(BatchNormalizatio...	128)		
conv3_block3_2_relu	(None, 16, 16,	0	
conv3_block3_2_b...			
(Activation)	128)		
conv3_block3_3_conv	(None, 16, 16,	66,048	
conv3_block3_2_r...			
(Conv2D)	512)		

conv3_block3_3_bn	(None, 16, 16,	2,048	
conv3_block3_3_c...	(BatchNormalizatio...	512)	
conv3_block3_add	(None, 16, 16,	0	
conv3_block2_out...	(Add)	512)	
conv3_block3_3_b...			
conv3_block3_out	(None, 16, 16,	0	
conv3_block3_add...	(Activation)	512)	
conv3_block4_1_conv	(None, 16, 16,	65,664	
conv3_block3_out...	(Conv2D)	128)	
conv3_block4_1_bn	(None, 16, 16,	512	
conv3_block4_1_c...	(BatchNormalizatio...	128)	
conv3_block4_1_relu	(None, 16, 16,	0	
conv3_block4_1_b...	(Activation)	128)	
conv3_block4_2_conv	(None, 16, 16,	147,584	
conv3_block4_1_r...	(Conv2D)	128)	
conv3_block4_2_bn	(None, 16, 16,	512	
conv3_block4_2_c...	(BatchNormalizatio...	128)	

conv3_block4_2_relu	(None, 16, 16,	0	
conv3_block4_2_bn	(Activation)	128)	
conv3_block4_3_conv	(None, 16, 16,	66,048	
conv3_block4_2_r...	(Conv2D)	512)	
conv3_block4_3_bn	(None, 16, 16,	2,048	
conv3_block4_3_c...	(BatchNormalizatio...	512)	
conv3_block4_add	(None, 16, 16,	0	
conv3_block3_out...	(Add)	512)	
conv3_block4_3_b...			
conv3_block4_out	(None, 16, 16,	0	
conv3_block4_add...	(Activation)	512)	
conv4_block1_1_conv	(None, 8, 8, 256)	131,328	
conv3_block4_out...	(Conv2D)		
conv4_block1_1_bn	(None, 8, 8, 256)	1,024	
conv4_block1_1_c...	(BatchNormalizatio...		
conv4_block1_1_relu	(None, 8, 8, 256)	0	
conv4_block1_1_b...	(Activation)		

conv4_block1_2_conv	(None, 8, 8, 256)	590,080	
conv4_block1_1_r... (Conv2D)			
conv4_block1_2_bn	(None, 8, 8, 256)	1,024	
conv4_block1_2_c... (BatchNormalizatio...			
conv4_block1_2_relu	(None, 8, 8, 256)	0	
conv4_block1_2_b... (Activation)			
conv4_block1_0_conv	(None, 8, 8,	525,312	
conv3_block4_out... (Conv2D)	1024)		
conv4_block1_3_conv	(None, 8, 8,	263,168	
conv4_block1_2_r... (Conv2D)	1024)		
conv4_block1_0_bn	(None, 8, 8,	4,096	
conv4_block1_0_c... (BatchNormalizatio...	1024)		
conv4_block1_3_bn	(None, 8, 8,	4,096	
conv4_block1_3_c... (BatchNormalizatio...	1024)		
conv4_block1_add	(None, 8, 8,	0	
conv4_block1_0_b... (Add)	1024)		
conv4_block1_3_b...			
conv4_block1_out	(None, 8, 8,	0	

conv4_block1_add...	(Activation)	1024)		
conv4_block2_1_conv	(None, 8, 8, 256)		262,400	
conv4_block1_out...	(Conv2D)			
conv4_block2_1_bn	(None, 8, 8, 256)		1,024	
conv4_block2_1_c...	(BatchNormalizatio...			
conv4_block2_1_relu	(None, 8, 8, 256)		0	
conv4_block2_1_b...	(Activation)			
conv4_block2_2_conv	(None, 8, 8, 256)		590,080	
conv4_block2_1_r...	(Conv2D)			
conv4_block2_2_bn	(None, 8, 8, 256)		1,024	
conv4_block2_2_c...	(BatchNormalizatio...			
conv4_block2_2_relu	(None, 8, 8, 256)		0	
conv4_block2_2_b...	(Activation)			
conv4_block2_3_conv	(None, 8, 8,		263,168	
conv4_block2_2_r...	(Conv2D)	1024)		
conv4_block2_3_bn	(None, 8, 8,		4,096	
conv4_block2_3_c...				

(BatchNormalization)	1024)		
conv4_block2_add conv4_block1_out...	(None, 8, 8,	0	
(Add)	1024)		
conv4_block2_3_b...			
conv4_block2_out conv4_block2_add...	(None, 8, 8,	0	
(Activation)	1024)		
conv4_block3_1_conv conv4_block2_out...	(None, 8, 8, 256)	262,400	
(Conv2D)			
conv4_block3_1_bn conv4_block3_1_c...	(None, 8, 8, 256)	1,024	
(BatchNormalization)			
conv4_block3_1_relu conv4_block3_1_b...	(None, 8, 8, 256)	0	
(Activation)			
conv4_block3_2_conv conv4_block3_1_r...	(None, 8, 8, 256)	590,080	
(Conv2D)			
conv4_block3_2_bn conv4_block3_2_c...	(None, 8, 8, 256)	1,024	
(BatchNormalization)			
conv4_block3_2_relu conv4_block3_2_b...	(None, 8, 8, 256)	0	
(Activation)			

conv4_block3_3_conv	(None, 8, 8,	263,168	
conv4_block3_2_r... (Conv2D)	1024)		
conv4_block3_3_bn	(None, 8, 8,	4,096	
conv4_block3_3_c... (BatchNormalizatio...	1024)		
conv4_block3_add	(None, 8, 8,	0	
conv4_block2_out... (Add)	1024)		
conv4_block3_3_b...			
conv4_block3_out	(None, 8, 8,	0	
conv4_block3_add... (Activation)	1024)		
conv4_block4_1_conv	(None, 8, 8, 256)	262,400	
conv4_block3_out... (Conv2D)			
conv4_block4_1_bn	(None, 8, 8, 256)	1,024	
conv4_block4_1_c... (BatchNormalizatio...			
conv4_block4_1_relu	(None, 8, 8, 256)	0	
conv4_block4_1_b... (Activation)			
conv4_block4_2_conv	(None, 8, 8, 256)	590,080	
conv4_block4_1_r... (Conv2D)			

conv4_block4_2_bn	(None, 8, 8, 256)	1,024	
conv4_block4_2_c...	(BatchNormalizatio...		
conv4_block4_2_relu	(None, 8, 8, 256)	0	
conv4_block4_2_b...	(Activation)		
conv4_block4_3_conv	(None, 8, 8,	263,168	
conv4_block4_2_r...	(Conv2D)	1024)	
conv4_block4_3_bn	(None, 8, 8,	4,096	
conv4_block4_3_c...	(BatchNormalizatio...	1024)	
conv4_block4_add	(None, 8, 8,	0	
conv4_block3_out...	(Add)	1024)	
conv4_block4_3_b...			
conv4_block4_out	(None, 8, 8,	0	
conv4_block4_add...	(Activation)	1024)	
conv4_block5_1_conv	(None, 8, 8, 256)	262,400	
conv4_block4_out...	(Conv2D)		
conv4_block5_1_bn	(None, 8, 8, 256)	1,024	
conv4_block5_1_c...	(BatchNormalizatio...		

conv4_block5_1_relu	(None, 8, 8, 256)	0	
conv4_block5_1_bn			
(Activation)			
conv4_block5_2_conv	(None, 8, 8, 256)	590,080	
conv4_block5_1_relu			
(Conv2D)			
conv4_block5_2_bn	(None, 8, 8, 256)	1,024	
conv4_block5_2_conv			
(BatchNormalization)			
conv4_block5_2_relu	(None, 8, 8, 256)	0	
conv4_block5_2_bn			
(Activation)			
conv4_block5_3_conv	(None, 8, 8,	263,168	
conv4_block5_2_relu			
(Conv2D)	1024)		
conv4_block5_3_bn	(None, 8, 8,	4,096	
conv4_block5_3_conv			
(BatchNormalization)	1024)		
conv4_block5_add	(None, 8, 8,	0	
conv4_block4_out			
(Add)	1024)		
conv4_block5_3_bn			
conv4_block5_out	(None, 8, 8,	0	
conv4_block5_add			
(Activation)	1024)		

conv4_block6_1_conv	(None, 8, 8, 256)	262,400	
conv4_block5_out...			
(Conv2D)			
conv4_block6_1_bn	(None, 8, 8, 256)	1,024	
conv4_block6_1_c...			
(BatchNormalizatio...			
conv4_block6_1_relu	(None, 8, 8, 256)	0	
conv4_block6_1_b...			
(Activation)			
conv4_block6_2_conv	(None, 8, 8, 256)	590,080	
conv4_block6_1_r...			
(Conv2D)			
conv4_block6_2_bn	(None, 8, 8, 256)	1,024	
conv4_block6_2_c...			
(BatchNormalizatio...			
conv4_block6_2_relu	(None, 8, 8, 256)	0	
conv4_block6_2_b...			
(Activation)			
conv4_block6_3_conv	(None, 8, 8,	263,168	
conv4_block6_2_r...			
(Conv2D)	1024)		
conv4_block6_3_bn	(None, 8, 8,	4,096	
conv4_block6_3_c...			
(BatchNormalizatio...	1024)		
conv4_block6_add	(None, 8, 8,	0	

conv4_block5_out...				
(Add)		1024)		
conv4_block6_3_b...				
<hr/>				
conv4_block6_out		(None, 8, 8,	0	
conv4_block6_add...		(Activation)	1024)	
<hr/>				
conv5_block1_1_conv		(None, 4, 4, 512)	524,800	
conv4_block6_out...		(Conv2D)		
<hr/>				
conv5_block1_1_bn		(None, 4, 4, 512)	2,048	
conv5_block1_1_c...		(BatchNormalizatio...		
<hr/>				
conv5_block1_1_relu		(None, 4, 4, 512)	0	
conv5_block1_1_b...		(Activation)		
<hr/>				
conv5_block1_2_conv		(None, 4, 4, 512)	2,359,808	
conv5_block1_1_r...		(Conv2D)		
<hr/>				
conv5_block1_2_bn		(None, 4, 4, 512)	2,048	
conv5_block1_2_c...		(BatchNormalizatio...		
<hr/>				
conv5_block1_2_relu		(None, 4, 4, 512)	0	
conv5_block1_2_b...		(Activation)		
<hr/>				
conv5_block1_0_conv		(None, 4, 4,	2,099,200	
conv4_block6_out...				

(Conv2D)	2048)		
conv5_block1_3_conv	(None, 4, 4,	1,050,624	
conv5_block1_2_r...	(Conv2D)	2048)	
conv5_block1_0_bn	(None, 4, 4,	8,192	
conv5_block1_0_c...	(BatchNormalizatio...	2048)	
conv5_block1_3_bn	(None, 4, 4,	8,192	
conv5_block1_3_c...	(BatchNormalizatio...	2048)	
conv5_block1_add	(None, 4, 4,	0	
conv5_block1_0_b...	(Add)	2048)	
conv5_block1_3_b...			
conv5_block1_out	(None, 4, 4,	0	
conv5_block1_add...	(Activation)	2048)	
conv5_block2_1_conv	(None, 4, 4, 512)	1,049,088	
conv5_block1_out...	(Conv2D)		
conv5_block2_1_bn	(None, 4, 4, 512)	2,048	
conv5_block2_1_c...	(BatchNormalizatio...		
conv5_block2_1_relu	(None, 4, 4, 512)	0	
conv5_block2_1_b...	(Activation)		

conv5_block2_2_conv	(None, 4, 4, 512)	2,359,808	
conv5_block2_1_r... (Conv2D)			
conv5_block2_2_bn	(None, 4, 4, 512)	2,048	
conv5_block2_2_c... (BatchNormalizatio...			
conv5_block2_2_relu	(None, 4, 4, 512)	0	
conv5_block2_2_b... (Activation)			
conv5_block2_3_conv	(None, 4, 4,	1,050,624	
conv5_block2_2_r... (Conv2D)	2048)		
conv5_block2_3_bn	(None, 4, 4,	8,192	
conv5_block2_3_c... (BatchNormalizatio...	2048)		
conv5_block2_add	(None, 4, 4,	0	
conv5_block1_out... (Add)	2048)		
conv5_block2_3_b...			
conv5_block2_out	(None, 4, 4,	0	
conv5_block2_add... (Activation)	2048)		
conv5_block3_1_conv	(None, 4, 4, 512)	1,049,088	
conv5_block2_out... (Conv2D)			

conv5_block3_1_bn	(None, 4, 4, 512)	2,048	
conv5_block3_1_c...	(BatchNormalizatio...		
conv5_block3_1_relu	(None, 4, 4, 512)	0	
conv5_block3_1_b...	(Activation)		
conv5_block3_2_conv	(None, 4, 4, 512)	2,359,808	
conv5_block3_1_r...	(Conv2D)		
conv5_block3_2_bn	(None, 4, 4, 512)	2,048	
conv5_block3_2_c...	(BatchNormalizatio...		
conv5_block3_2_relu	(None, 4, 4, 512)	0	
conv5_block3_2_b...	(Activation)		
conv5_block3_3_conv	(None, 4, 4,	1,050,624	
conv5_block3_2_r...	(Conv2D)	2048)	
conv5_block3_3_bn	(None, 4, 4,	8,192	
conv5_block3_3_c...	(BatchNormalizatio...	2048)	
conv5_block3_add	(None, 4, 4,	0	
conv5_block2_out...	(Add)	2048)	
conv5_block3_3_b...			

conv5_block3_out	(None, 4, 4,	0	
conv5_block3_add...	(Activation)	2048)	

Total params: 23,587,712 (89.98 MB)

Trainable params: 0 (0.00 B)

Non-trainable params: 23,587,712 (89.98 MB)

4. Add Custom Layers

- Add Global Average Pooling, Dropout, Dense layers, and a final softmax output layer.

```
from tensorflow.keras.layers import GlobalAveragePooling2D, Dropout, Dense
from tensorflow.keras.models import Model
```

```
# Add custom layers on top of the base_model
x = GlobalAveragePooling2D()(base_model.output)
x = Dropout(0.5)(x)
x = Dense(128, activation='relu')(x)
output = Dense(len(classes), activation='softmax')(x)
```

```
model = Model(inputs=base_model.input, outputs=output)
model.summary()
```

Model: "functional_1"

Layer (type)	Output Shape	Param #	Connected to
input_layer_1 (InputLayer)	(None, 128, 128, 3)	0	-
conv1_pad input_layer_1[0]...	(None, 134, 134, 3)	0	
(ZeroPadding2D)			
conv1_conv (Conv2D)	(None, 64, 64,	9,472	conv1_pad[0]

[0]		64)		
	conv1_bn	(None, 64, 64,	256	conv1_conv[0]
[0]	(BatchNormalizatio...	64)		
	conv1_relu	(None, 64, 64,	0	conv1_bn[0]
[0]	(Activation)	64)		
	pool1_pad	(None, 66, 66,	0	conv1_relu[0]
[0]	(ZeroPadding2D)	64)		
	pool1_pool	(None, 32, 32,	0	pool1_pad[0]
[0]	(MaxPooling2D)	64)		
	conv2_block1_1_conv	(None, 32, 32,	4,160	pool1_pool[0]
[0]	(Conv2D)	64)		
	conv2_block1_1_bn	(None, 32, 32,	256	
	conv2_block1_1_c...	(BatchNormalizatio...	64)	
	conv2_block1_1_relu	(None, 32, 32,	0	
	conv2_block1_1_b...	(Activation)	64)	
	conv2_block1_2_conv	(None, 32, 32,	36,928	
	conv2_block1_1_r...			

(Conv2D)	64)		
conv2_block1_2_bn	(None, 32, 32,	256	
conv2_block1_2_c...	(BatchNormalizatio...	64)	
conv2_block1_2_relu	(None, 32, 32,	0	
conv2_block1_2_b...	(Activation)	64)	
conv2_block1_0_conv	(None, 32, 32,	16,640	pool1_pool[0]
[0]	(Conv2D)	256)	
conv2_block1_3_conv	(None, 32, 32,	16,640	
conv2_block1_2_r...	(Conv2D)	256)	
conv2_block1_0_bn	(None, 32, 32,	1,024	
conv2_block1_0_c...	(BatchNormalizatio...	256)	
conv2_block1_3_bn	(None, 32, 32,	1,024	
conv2_block1_3_c...	(BatchNormalizatio...	256)	
conv2_block1_add	(None, 32, 32,	0	
conv2_block1_0_b...	(Add)	256)	
conv2_block1_3_b...			
conv2_block1_out	(None, 32, 32,	0	
conv2_block1_add...	(Activation)	256)	

conv2_block2_1_conv conv2_block1_out...	(None, 32, 32, 64)	16,448	
conv2_block2_1_bn conv2_block2_1_c...	(None, 32, 32, 64)	256	
conv2_block2_1_relu conv2_block2_1_b...	(None, 32, 32, 64)	0	
conv2_block2_2_conv conv2_block2_1_r...	(None, 32, 32, 64)	36,928	
conv2_block2_2_bn conv2_block2_2_c...	(None, 32, 32, 64)	256	
conv2_block2_2_relu conv2_block2_2_b...	(None, 32, 32, 64)	0	
conv2_block2_3_conv conv2_block2_2_r...	(None, 32, 32, 256)	16,640	
conv2_block2_3_bn conv2_block2_3_c...	(None, 32, 32, 256)	1,024	

conv2_block2_add conv2_block1_out...	(None, 32, 32,	0	
(Add)	256)		
conv2_block2_3_b...			
conv2_block2_out conv2_block2_add...	(None, 32, 32,	0	
(Activation)	256)		
conv2_block3_1_conv conv2_block2_out...	(None, 32, 32,	16,448	
(Conv2D)	64)		
conv2_block3_1_bn conv2_block3_1_c...	(None, 32, 32,	256	
(BatchNormalizatio...	64)		
conv2_block3_1_relu conv2_block3_1_b...	(None, 32, 32,	0	
(Activation)	64)		
conv2_block3_2_conv conv2_block3_1_r...	(None, 32, 32,	36,928	
(Conv2D)	64)		
conv2_block3_2_bn conv2_block3_2_c...	(None, 32, 32,	256	
(BatchNormalizatio...	64)		
conv2_block3_2_relu conv2_block3_2_b...	(None, 32, 32,	0	
(Activation)	64)		

conv2_block3_3_conv	(None, 32, 32,	16,640
conv2_block3_2_r...	(Conv2D)	256)
conv2_block3_3_bn	(None, 32, 32,	1,024
conv2_block3_3_c...	(BatchNormalizatio...	256)
conv2_block3_add	(None, 32, 32,	0
conv2_block2_out...	(Add)	256)
conv2_block3_3_b...		
conv2_block3_out	(None, 32, 32,	0
conv2_block3_add...	(Activation)	256)
conv3_block1_1_conv	(None, 16, 16,	32,896
conv2_block3_out...	(Conv2D)	128)
conv3_block1_1_bn	(None, 16, 16,	512
conv3_block1_1_c...	(BatchNormalizatio...	128)
conv3_block1_1_relu	(None, 16, 16,	0
conv3_block1_1_b...	(Activation)	128)
conv3_block1_2_conv	(None, 16, 16,	147,584
conv3_block1_1_r...	(Conv2D)	128)

conv3_block1_2_bn	(None, 16, 16,	512	
conv3_block1_2_c...	(BatchNormalizatio...	128)	
conv3_block1_2_relu	(None, 16, 16,	0	
conv3_block1_2_b...	(Activation)	128)	
conv3_block1_0_conv	(None, 16, 16,	131,584	
conv2_block3_out...	(Conv2D)	512)	
conv3_block1_3_conv	(None, 16, 16,	66,048	
conv3_block1_2_r...	(Conv2D)	512)	
conv3_block1_0_bn	(None, 16, 16,	2,048	
conv3_block1_0_c...	(BatchNormalizatio...	512)	
conv3_block1_3_bn	(None, 16, 16,	2,048	
conv3_block1_3_c...	(BatchNormalizatio...	512)	
conv3_block1_add	(None, 16, 16,	0	
conv3_block1_0_b...	(Add)	512)	
conv3_block1_3_b...			
conv3_block1_out	(None, 16, 16,	0	
conv3_block1_add...	(Activation)	512)	
conv3_block2_1_conv	(None, 16, 16,	65,664	

conv3_block1_out...	(Conv2D)	128)		
conv3_block2_1_bn	(None, 16, 16,	512		
conv3_block2_1_c...	(BatchNormalizatio...	128)		
conv3_block2_1_relu	(None, 16, 16,	0		
conv3_block2_1_b...	(Activation)	128)		
conv3_block2_2_conv	(None, 16, 16,	147,584		
conv3_block2_1_r...	(Conv2D)	128)		
conv3_block2_2_bn	(None, 16, 16,	512		
conv3_block2_2_c...	(BatchNormalizatio...	128)		
conv3_block2_2_relu	(None, 16, 16,	0		
conv3_block2_2_b...	(Activation)	128)		
conv3_block2_3_conv	(None, 16, 16,	66,048		
conv3_block2_2_r...	(Conv2D)	512)		
conv3_block2_3_bn	(None, 16, 16,	2,048		
conv3_block2_3_c...	(BatchNormalizatio...	512)		
conv3_block2_add	(None, 16, 16,	0		
conv3_block1_out...				

(Add)	512)		
conv3_block2_3_b...			
conv3_block2_out	(None, 16, 16,	0	
conv3_block2_add...			
(Activation)	512)		
conv3_block3_1_conv	(None, 16, 16,	65,664	
conv3_block2_out...			
(Conv2D)	128)		
conv3_block3_1_bn	(None, 16, 16,	512	
conv3_block3_1_c...			
(BatchNormalizatio...	128)		
conv3_block3_1_relu	(None, 16, 16,	0	
conv3_block3_1_b...			
(Activation)	128)		
conv3_block3_2_conv	(None, 16, 16,	147,584	
conv3_block3_1_r...			
(Conv2D)	128)		
conv3_block3_2_bn	(None, 16, 16,	512	
conv3_block3_2_c...			
(BatchNormalizatio...	128)		
conv3_block3_2_relu	(None, 16, 16,	0	
conv3_block3_2_b...			
(Activation)	128)		
conv3_block3_3_conv	(None, 16, 16,	66,048	
conv3_block3_2_r...			
(Conv2D)	512)		

conv3_block3_3_bn	(None, 16, 16,	2,048	
conv3_block3_3_c...	(BatchNormalizatio...	512)	
conv3_block3_add	(None, 16, 16,	0	
conv3_block2_out...	(Add)	512)	
conv3_block3_3_b...			
conv3_block3_out	(None, 16, 16,	0	
conv3_block3_add...	(Activation)	512)	
conv3_block4_1_conv	(None, 16, 16,	65,664	
conv3_block3_out...	(Conv2D)	128)	
conv3_block4_1_bn	(None, 16, 16,	512	
conv3_block4_1_c...	(BatchNormalizatio...	128)	
conv3_block4_1_relu	(None, 16, 16,	0	
conv3_block4_1_b...	(Activation)	128)	
conv3_block4_2_conv	(None, 16, 16,	147,584	
conv3_block4_1_r...	(Conv2D)	128)	
conv3_block4_2_bn	(None, 16, 16,	512	
conv3_block4_2_c...	(BatchNormalizatio...	128)	

conv3_block4_2_relu	(None, 16, 16,	0	
conv3_block4_2_bn	(Activation)	128)	
conv3_block4_3_conv	(None, 16, 16,	66,048	
conv3_block4_2_r...	(Conv2D)	512)	
conv3_block4_3_bn	(None, 16, 16,	2,048	
conv3_block4_3_c...	(BatchNormalizatio...	512)	
conv3_block4_add	(None, 16, 16,	0	
conv3_block3_out...	(Add)	512)	
conv3_block4_3_b...			
conv3_block4_out	(None, 16, 16,	0	
conv3_block4_add...	(Activation)	512)	
conv4_block1_1_conv	(None, 8, 8, 256)	131,328	
conv3_block4_out...	(Conv2D)		
conv4_block1_1_bn	(None, 8, 8, 256)	1,024	
conv4_block1_1_c...	(BatchNormalizatio...		
conv4_block1_1_relu	(None, 8, 8, 256)	0	
conv4_block1_1_b...	(Activation)		

conv4_block1_2_conv	(None, 8, 8, 256)	590,080	
conv4_block1_1_r... (Conv2D)			
conv4_block1_2_bn	(None, 8, 8, 256)	1,024	
conv4_block1_2_c... (BatchNormalizatio...			
conv4_block1_2_relu	(None, 8, 8, 256)	0	
conv4_block1_2_b... (Activation)			
conv4_block1_0_conv	(None, 8, 8,	525,312	
conv3_block4_out... (Conv2D)	1024)		
conv4_block1_3_conv	(None, 8, 8,	263,168	
conv4_block1_2_r... (Conv2D)	1024)		
conv4_block1_0_bn	(None, 8, 8,	4,096	
conv4_block1_0_c... (BatchNormalizatio...	1024)		
conv4_block1_3_bn	(None, 8, 8,	4,096	
conv4_block1_3_c... (BatchNormalizatio...	1024)		
conv4_block1_add	(None, 8, 8,	0	
conv4_block1_0_b... (Add)	1024)		
conv4_block1_3_b...			
conv4_block1_out	(None, 8, 8,	0	

conv4_block1_add...	(Activation)	1024)		
conv4_block2_1_conv	(None, 8, 8, 256)		262,400	
conv4_block1_out...	(Conv2D)			
conv4_block2_1_bn	(None, 8, 8, 256)		1,024	
conv4_block2_1_c...	(BatchNormalizatio...			
conv4_block2_1_relu	(None, 8, 8, 256)		0	
conv4_block2_1_b...	(Activation)			
conv4_block2_2_conv	(None, 8, 8, 256)		590,080	
conv4_block2_1_r...	(Conv2D)			
conv4_block2_2_bn	(None, 8, 8, 256)		1,024	
conv4_block2_2_c...	(BatchNormalizatio...			
conv4_block2_2_relu	(None, 8, 8, 256)		0	
conv4_block2_2_b...	(Activation)			
conv4_block2_3_conv	(None, 8, 8,		263,168	
conv4_block2_2_r...	(Conv2D)	1024)		
conv4_block2_3_bn	(None, 8, 8,		4,096	
conv4_block2_3_c...				

(BatchNormalization)	1024)		
conv4_block2_add conv4_block1_out...	(None, 8, 8,	0	
(Add)	1024)		
conv4_block2_3_b...			
conv4_block2_out conv4_block2_add...	(None, 8, 8,	0	
(Activation)	1024)		
conv4_block3_1_conv conv4_block2_out...	(None, 8, 8, 256)	262,400	
(Conv2D)			
conv4_block3_1_bn conv4_block3_1_c...	(None, 8, 8, 256)	1,024	
(BatchNormalization)			
conv4_block3_1_relu conv4_block3_1_b...	(None, 8, 8, 256)	0	
(Activation)			
conv4_block3_2_conv conv4_block3_1_r...	(None, 8, 8, 256)	590,080	
(Conv2D)			
conv4_block3_2_bn conv4_block3_2_c...	(None, 8, 8, 256)	1,024	
(BatchNormalization)			
conv4_block3_2_relu conv4_block3_2_b...	(None, 8, 8, 256)	0	
(Activation)			

conv4_block3_3_conv	(None, 8, 8,	263,168	
conv4_block3_2_r...	(Conv2D)	1024)	
conv4_block3_3_bn	(None, 8, 8,	4,096	
conv4_block3_3_c...	(BatchNormalizatio...	1024)	
conv4_block3_add	(None, 8, 8,	0	
conv4_block2_out...	(Add)	1024)	
conv4_block3_3_b...			
conv4_block3_out	(None, 8, 8,	0	
conv4_block3_add...	(Activation)	1024)	
conv4_block4_1_conv	(None, 8, 8, 256)	262,400	
conv4_block3_out...	(Conv2D)		
conv4_block4_1_bn	(None, 8, 8, 256)	1,024	
conv4_block4_1_c...	(BatchNormalizatio...		
conv4_block4_1_relu	(None, 8, 8, 256)	0	
conv4_block4_1_b...	(Activation)		
conv4_block4_2_conv	(None, 8, 8, 256)	590,080	
conv4_block4_1_r...	(Conv2D)		

conv4_block4_2_bn	(None, 8, 8, 256)	1,024	
conv4_block4_2_c...	(BatchNormalizatio...		
conv4_block4_2_relu	(None, 8, 8, 256)	0	
conv4_block4_2_b...	(Activation)		
conv4_block4_3_conv	(None, 8, 8,	263,168	
conv4_block4_2_r...	(Conv2D)	1024)	
conv4_block4_3_bn	(None, 8, 8,	4,096	
conv4_block4_3_c...	(BatchNormalizatio...	1024)	
conv4_block4_add	(None, 8, 8,	0	
conv4_block3_out...	(Add)	1024)	
conv4_block4_3_b...			
conv4_block4_out	(None, 8, 8,	0	
conv4_block4_add...	(Activation)	1024)	
conv4_block5_1_conv	(None, 8, 8, 256)	262,400	
conv4_block4_out...	(Conv2D)		
conv4_block5_1_bn	(None, 8, 8, 256)	1,024	
conv4_block5_1_c...	(BatchNormalizatio...		

conv4_block5_1_relu	(None, 8, 8, 256)	0	
conv4_block5_1_bn			
(Activation)			
conv4_block5_2_conv	(None, 8, 8, 256)	590,080	
conv4_block5_1_relu			
(Conv2D)			
conv4_block5_2_bn	(None, 8, 8, 256)	1,024	
conv4_block5_2_conv			
(BatchNormalization)			
conv4_block5_2_relu	(None, 8, 8, 256)	0	
conv4_block5_2_bn			
(Activation)			
conv4_block5_3_conv	(None, 8, 8,	263,168	
conv4_block5_2_relu			
(Conv2D)	1024)		
conv4_block5_3_bn	(None, 8, 8,	4,096	
conv4_block5_3_conv			
(BatchNormalization)	1024)		
conv4_block5_add	(None, 8, 8,	0	
conv4_block4_out			
(Add)	1024)		
conv4_block5_3_bn			
conv4_block5_out	(None, 8, 8,	0	
conv4_block5_add			
(Activation)	1024)		

conv4_block6_1_conv	(None, 8, 8, 256)	262,400	
conv4_block5_out...			
(Conv2D)			
<hr/>			
conv4_block6_1_bn	(None, 8, 8, 256)	1,024	
conv4_block6_1_c...			
(BatchNormalizatio...			
<hr/>			
conv4_block6_1_relu	(None, 8, 8, 256)	0	
conv4_block6_1_b...			
(Activation)			
<hr/>			
conv4_block6_2_conv	(None, 8, 8, 256)	590,080	
conv4_block6_1_r...			
(Conv2D)			
<hr/>			
conv4_block6_2_bn	(None, 8, 8, 256)	1,024	
conv4_block6_2_c...			
(BatchNormalizatio...			
<hr/>			
conv4_block6_2_relu	(None, 8, 8, 256)	0	
conv4_block6_2_b...			
(Activation)			
<hr/>			
conv4_block6_3_conv	(None, 8, 8,	263,168	
conv4_block6_2_r...			
(Conv2D)	1024)		
<hr/>			
conv4_block6_3_bn	(None, 8, 8,	4,096	
conv4_block6_3_c...			
(BatchNormalizatio...	1024)		
<hr/>			
conv4_block6_add	(None, 8, 8,	0	

conv4_block5_out...				
(Add)		1024)		
conv4_block6_3_b...				
<hr/>				
conv4_block6_out		(None, 8, 8,	0	
conv4_block6_add...		(Activation)	1024)	
<hr/>				
conv5_block1_1_conv		(None, 4, 4, 512)	524,800	
conv4_block6_out...		(Conv2D)		
<hr/>				
conv5_block1_1_bn		(None, 4, 4, 512)	2,048	
conv5_block1_1_c...		(BatchNormalizatio...		
<hr/>				
conv5_block1_1_relu		(None, 4, 4, 512)	0	
conv5_block1_1_b...		(Activation)		
<hr/>				
conv5_block1_2_conv		(None, 4, 4, 512)	2,359,808	
conv5_block1_1_r...		(Conv2D)		
<hr/>				
conv5_block1_2_bn		(None, 4, 4, 512)	2,048	
conv5_block1_2_c...		(BatchNormalizatio...		
<hr/>				
conv5_block1_2_relu		(None, 4, 4, 512)	0	
conv5_block1_2_b...		(Activation)		
<hr/>				
conv5_block1_0_conv		(None, 4, 4,	2,099,200	
conv4_block6_out...				

(Conv2D)	2048)		
conv5_block1_3_conv	(None, 4, 4,	1,050,624	
conv5_block1_2_r...	(Conv2D)	2048)	
conv5_block1_0_bn	(None, 4, 4,	8,192	
conv5_block1_0_c...	(BatchNormalizatio...	2048)	
conv5_block1_3_bn	(None, 4, 4,	8,192	
conv5_block1_3_c...	(BatchNormalizatio...	2048)	
conv5_block1_add	(None, 4, 4,	0	
conv5_block1_0_b...	(Add)	2048)	
conv5_block1_3_b...			
conv5_block1_out	(None, 4, 4,	0	
conv5_block1_add...	(Activation)	2048)	
conv5_block2_1_conv	(None, 4, 4, 512)	1,049,088	
conv5_block1_out...	(Conv2D)		
conv5_block2_1_bn	(None, 4, 4, 512)	2,048	
conv5_block2_1_c...	(BatchNormalizatio...		
conv5_block2_1_relu	(None, 4, 4, 512)	0	
conv5_block2_1_b...	(Activation)		

conv5_block2_2_conv	(None, 4, 4, 512)	2,359,808	
conv5_block2_1_r... (Conv2D)			
conv5_block2_2_bn	(None, 4, 4, 512)	2,048	
conv5_block2_2_c... (BatchNormalizatio...			
conv5_block2_2_relu	(None, 4, 4, 512)	0	
conv5_block2_2_b... (Activation)			
conv5_block2_3_conv	(None, 4, 4,	1,050,624	
conv5_block2_2_r... (Conv2D)	2048)		
conv5_block2_3_bn	(None, 4, 4,	8,192	
conv5_block2_3_c... (BatchNormalizatio...	2048)		
conv5_block2_add	(None, 4, 4,	0	
conv5_block1_out... (Add)	2048)		
conv5_block2_3_b...			
conv5_block2_out	(None, 4, 4,	0	
conv5_block2_add... (Activation)	2048)		
conv5_block3_1_conv	(None, 4, 4, 512)	1,049,088	
conv5_block2_out... (Conv2D)			

conv5_block3_1_bn	(None, 4, 4, 512)	2,048	
conv5_block3_1_c...	(BatchNormalizatio...		
conv5_block3_1_relu	(None, 4, 4, 512)	0	
conv5_block3_1_b...	(Activation)		
conv5_block3_2_conv	(None, 4, 4, 512)	2,359,808	
conv5_block3_1_r...	(Conv2D)		
conv5_block3_2_bn	(None, 4, 4, 512)	2,048	
conv5_block3_2_c...	(BatchNormalizatio...		
conv5_block3_2_relu	(None, 4, 4, 512)	0	
conv5_block3_2_b...	(Activation)		
conv5_block3_3_conv	(None, 4, 4,	1,050,624	
conv5_block3_2_r...	(Conv2D)	2048)	
conv5_block3_3_bn	(None, 4, 4,	8,192	
conv5_block3_3_c...	(BatchNormalizatio...	2048)	
conv5_block3_add	(None, 4, 4,	0	
conv5_block2_out...	(Add)	2048)	
conv5_block3_3_b...			

conv5_block3_out conv5_block3_add...	(None, 4, 4,	0	
(Activation)	2048)		
<hr/>			
global_average_poo... conv5_block3_out...	(None, 2048)	0	
(GlobalAveragePool...			
<hr/>			
dropout_1 (Dropout)	(None, 2048)	0	
global_average_p...			
<hr/>			
dense_2 (Dense)	(None, 128)	262,272	dropout_1[0]
[0]			
<hr/>			
dense_3 (Dense)	(None, 3)	387	dense_2[0][0]
<hr/>			

Total params: 23,850,371 (90.98 MB)

Trainable params: 262,659 (1.00 MB)

Non-trainable params: 23,587,712 (89.98 MB)

5. Freeze Pre-trained Layers

- Freeze all layers in the ResNet base to train only the custom head.

```
for layer in base_model.layers:
    layer.trainable = False

# Optionally, check if all layers are frozen
print("All base_model layers frozen:", all(not layer.trainable for
layer in base_model.layers))
```

All base_model layers frozen: True

6. Compile and Train the Model (Initial Training)

- Use Adam optimizer, categorical crossentropy, and accuracy metrics.
- Add callbacks: EarlyStopping, ModelCheckpoint, ReduceLROnPlateau.

```
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint,
```


ReduceLR0nPlateau

Compile the model

```
model.compile(
    optimizer=Adam(learning_rate=0.0001),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)
```

Define callbacks

```
callbacks = [
    EarlyStopping(monitor='val_loss', patience=5,
        restore_best_weights=True, verbose=1),
    ModelCheckpoint('best_model.h5', monitor='val_loss',
        save_best_only=True, verbose=1),
    ReduceLR0nPlateau(monitor='val_loss', factor=0.5, patience=2,
        verbose=1)
]
```

Train the model

```
history = model.fit(
    train_gen,
    epochs=25,
    validation_data=val_gen,
    callbacks=callbacks
)
```

c:\Users\This PC\.conda\envs\infotech_course\Lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.

```
self._warn_if_super_not_called()
```

Epoch 1/25

66/66 ————— 0s 294ms/step - accuracy: 0.5577 - loss: 1.2418

Epoch 1: val_loss improved from inf to 0.32776, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 32s 392ms/step - accuracy: 0.5595 - loss: 1.2365 - val_accuracy: 0.8933 - val_loss: 0.3278 - learning_rate: 1.0000e-04

Epoch 2/25
66/66 ————— 0s 290ms/step - accuracy: 0.8269 - loss: 0.4896

Epoch 2: val_loss improved from 0.32776 to 0.23459, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 23s 355ms/step - accuracy: 0.8270 - loss: 0.4890 - val_accuracy: 0.9289 - val_loss: 0.2346 - learning_rate: 1.0000e-04

Epoch 3/25
66/66 ————— 0s 289ms/step - accuracy: 0.8841 - loss: 0.3163

Epoch 3: val_loss improved from 0.23459 to 0.21907, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 23s 354ms/step - accuracy: 0.8839 - loss: 0.3167 - val_accuracy: 0.9267 - val_loss: 0.2191 - learning_rate: 1.0000e-04

Epoch 4/25
66/66 ————— 0s 290ms/step - accuracy: 0.8729 - loss: 0.3434

Epoch 4: val_loss improved from 0.21907 to 0.21174, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 24s 356ms/step - accuracy: 0.8730 - loss: 0.3432 - val_accuracy: 0.9311 - val_loss: 0.2117 - learning_rate: 1.0000e-04

Epoch 5/25
66/66 ————— 0s 293ms/step - accuracy: 0.8906 - loss: 0.2978

Epoch 5: val_loss improved from 0.21174 to 0.20588, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via
`model.save()` or `keras.saving.save_model(model)`. This file format
is considered legacy. We recommend using instead the native Keras
format, e.g. `model.save('my_model.keras')` or
`keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 24s 359ms/step - accuracy: 0.8905 - loss:
0.2981 - val_accuracy: 0.9311 - val_loss: 0.2059 - learning_rate:
1.0000e-04

Epoch 6/25

66/66 ————— 0s 292ms/step - accuracy: 0.9103 - loss:
0.2478

Epoch 6: val_loss improved from 0.20588 to 0.19631, saving model to
best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via
`model.save()` or `keras.saving.save_model(model)`. This file format
is considered legacy. We recommend using instead the native Keras
format, e.g. `model.save('my_model.keras')` or
`keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 24s 358ms/step - accuracy: 0.9102 - loss:
0.2481 - val_accuracy: 0.9356 - val_loss: 0.1963 - learning_rate:
1.0000e-04

Epoch 7/25

66/66 ————— 0s 296ms/step - accuracy: 0.9057 - loss:
0.2496

Epoch 7: val_loss did not improve from 0.19631

66/66 ————— 24s 355ms/step - accuracy: 0.9057 - loss:
0.2497 - val_accuracy: 0.9356 - val_loss: 0.2035 - learning_rate:
1.0000e-04

Epoch 8/25

66/66 ————— 0s 301ms/step - accuracy: 0.9070 - loss:
0.2133

Epoch 8: val_loss did not improve from 0.19631

Epoch 8: ReduceLROnPlateau reducing learning rate to
4.999999873689376e-05.

66/66 ————— 24s 363ms/step - accuracy: 0.9071 - loss:
0.2133 - val_accuracy: 0.9356 - val_loss: 0.1990 - learning_rate:
1.0000e-04

Epoch 9/25

66/66 ————— 0s 295ms/step - accuracy: 0.9253 - loss:
0.1990

Epoch 9: val_loss improved from 0.19631 to 0.19373, saving model to
best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via
`model.save()` or `keras.saving.save_model(model)`. This file format
is considered legacy. We recommend using instead the native Keras

format, e.g. ``model.save('my_model.keras')`` or
``keras.saving.save_model(model, 'my_model.keras')``.

66/66 ————— 24s 363ms/step - accuracy: 0.9252 - loss:
0.1992 - val_accuracy: 0.9356 - val_loss: 0.1937 - learning_rate:
5.0000e-05

Epoch 10/25

66/66 ————— 0s 305s/step - accuracy: 0.9168 - loss:
0.2070

Epoch 10: val_loss improved from 0.19373 to 0.19241, saving model to
best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via
``model.save()`` or ``keras.saving.save_model(model)``. This file format
is considered legacy. We recommend using instead the native Keras
format, e.g. ``model.save('my_model.keras')`` or
``keras.saving.save_model(model, 'my_model.keras')``.

66/66 ————— 19826s 305s/step - accuracy: 0.9168 - loss:
0.2071 - val_accuracy: 0.9400 - val_loss: 0.1924 - learning_rate:
5.0000e-05

Epoch 11/25

66/66 ————— 0s 309ms/step - accuracy: 0.9218 - loss:
0.2028

Epoch 11: val_loss improved from 0.19241 to 0.18791, saving model to
best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via
``model.save()`` or ``keras.saving.save_model(model)``. This file format
is considered legacy. We recommend using instead the native Keras
format, e.g. ``model.save('my_model.keras')`` or
``keras.saving.save_model(model, 'my_model.keras')``.

66/66 ————— 25s 380ms/step - accuracy: 0.9217 - loss:
0.2030 - val_accuracy: 0.9400 - val_loss: 0.1879 - learning_rate:
5.0000e-05

Epoch 12/25

66/66 ————— 0s 317ms/step - accuracy: 0.9033 - loss:
0.2367

Epoch 12: val_loss did not improve from 0.18791

66/66 ————— 25s 385ms/step - accuracy: 0.9034 - loss:
0.2364 - val_accuracy: 0.9422 - val_loss: 0.1887 - learning_rate:
5.0000e-05

Epoch 13/25

66/66 ————— 0s 340ms/step - accuracy: 0.9207 - loss:
0.2116

Epoch 13: val_loss did not improve from 0.18791

Epoch 13: ReduceLROnPlateau reducing learning rate to
2.499999936844688e-05.

66/66 ————— 27s 406ms/step - accuracy: 0.9208 - loss: 0.2115 - val_accuracy: 0.9422 - val_loss: 0.1900 - learning_rate: 5.0000e-05

Epoch 14/25

66/66 ————— 0s 336ms/step - accuracy: 0.9239 - loss: 0.2007

Epoch 14: val_loss did not improve from 0.18791

66/66 ————— 26s 400ms/step - accuracy: 0.9239 - loss: 0.2007 - val_accuracy: 0.9422 - val_loss: 0.1891 - learning_rate: 2.5000e-05

Epoch 15/25

66/66 ————— 0s 316ms/step - accuracy: 0.9188 - loss: 0.2093

Epoch 15: val_loss did not improve from 0.18791

Epoch 15: ReduceLROnPlateau reducing learning rate to 1.249999968422344e-05.

66/66 ————— 25s 378ms/step - accuracy: 0.9189 - loss: 0.2091 - val_accuracy: 0.9400 - val_loss: 0.1885 - learning_rate: 2.5000e-05

Epoch 16/25

66/66 ————— 0s 302ms/step - accuracy: 0.9098 - loss: 0.2184

Epoch 16: val_loss improved from 0.18791 to 0.18632, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 24s 369ms/step - accuracy: 0.9098 - loss: 0.2183 - val_accuracy: 0.9444 - val_loss: 0.1863 - learning_rate: 1.2500e-05

Epoch 17/25

66/66 ————— 0s 299ms/step - accuracy: 0.9153 - loss: 0.2165

Epoch 17: val_loss improved from 0.18632 to 0.18520, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 24s 367ms/step - accuracy: 0.9154 - loss: 0.2163 - val_accuracy: 0.9444 - val_loss: 0.1852 - learning_rate: 1.2500e-05

```

Epoch 18/25
66/66 _____ 0s 299ms/step - accuracy: 0.9325 - loss:
0.1946
Epoch 18: val_loss did not improve from 0.18520
66/66 _____ 24s 361ms/step - accuracy: 0.9326 - loss:
0.1944 - val_accuracy: 0.9400 - val_loss: 0.1870 - learning_rate:
1.2500e-05
Epoch 19/25
66/66 _____ 0s 300ms/step - accuracy: 0.9262 - loss:
0.1922
Epoch 19: val_loss did not improve from 0.18520

Epoch 19: ReduceLROnPlateau reducing learning rate to
6.24999984211172e-06.
66/66 _____ 24s 361ms/step - accuracy: 0.9262 - loss:
0.1923 - val_accuracy: 0.9422 - val_loss: 0.1862 - learning_rate:
1.2500e-05
Epoch 20/25
66/66 _____ 0s 312ms/step - accuracy: 0.9190 - loss:
0.2148
Epoch 20: val_loss did not improve from 0.18520
66/66 _____ 25s 386ms/step - accuracy: 0.9191 - loss:
0.2145 - val_accuracy: 0.9400 - val_loss: 0.1860 - learning_rate:
6.2500e-06
Epoch 21/25
66/66 _____ 0s 317ms/step - accuracy: 0.9116 - loss:
0.2110
Epoch 21: val_loss did not improve from 0.18520

Epoch 21: ReduceLROnPlateau reducing learning rate to
3.12499992105586e-06.
66/66 _____ 25s 379ms/step - accuracy: 0.9117 - loss:
0.2109 - val_accuracy: 0.9400 - val_loss: 0.1862 - learning_rate:
6.2500e-06
Epoch 22/25
66/66 _____ 0s 314ms/step - accuracy: 0.9223 - loss:
0.1886
Epoch 22: val_loss did not improve from 0.18520
66/66 _____ 25s 374ms/step - accuracy: 0.9222 - loss:
0.1888 - val_accuracy: 0.9422 - val_loss: 0.1857 - learning_rate:
3.1250e-06
Epoch 22: early stopping
Restoring model weights from the end of the best epoch: 17.

```

7. Fine-tune the Model

- Unfreeze some top layers of ResNet for fine-tuning.
- Retrain with a lower learning rate.

```

# Unfreeze the top layers of the base_model for fine-tuning
for layer in base_model.layers[-10:]:
    layer.trainable = True

# Recompile the model with a lower learning rate
model.compile(
    optimizer=Adam(learning_rate=1e-6),
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

# Retrain the model
fine_tune_history = model.fit(
    train_gen,
    epochs=10,
    validation_data=val_gen,
    callbacks=callbacks
)

Epoch 1/10
66/66 _____ 0s 339ms/step - accuracy: 0.9268 - loss:
0.1968
Epoch 1: val_loss improved from 0.18520 to 0.17922, saving model to
best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via
`model.save()` or `keras.saving.save_model(model)`. This file format
is considered legacy. We recommend using instead the native Keras
format, e.g. `model.save('my_model.keras')` or
`keras.saving.save_model(model, 'my_model.keras')`.

66/66 _____ 37s 444ms/step - accuracy: 0.9268 - loss:
0.1968 - val_accuracy: 0.9400 - val_loss: 0.1792 - learning_rate:
1.0000e-06
Epoch 2/10
66/66 _____ 0s 358ms/step - accuracy: 0.9181 - loss:
0.2001
Epoch 2: val_loss improved from 0.17922 to 0.17822, saving model to
best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via
`model.save()` or `keras.saving.save_model(model)`. This file format
is considered legacy. We recommend using instead the native Keras
format, e.g. `model.save('my_model.keras')` or
`keras.saving.save_model(model, 'my_model.keras')`.

66/66 _____ 29s 432ms/step - accuracy: 0.9181 - loss:
0.2001 - val_accuracy: 0.9378 - val_loss: 0.1782 - learning_rate:
1.0000e-06
Epoch 3/10
66/66 _____ 0s 342ms/step - accuracy: 0.9096 - loss:

```

0.2156

Epoch 3: val_loss improved from 0.17822 to 0.17786, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 27s 409ms/step - accuracy: 0.9096 - loss: 0.2154 - val_accuracy: 0.9378 - val_loss: 0.1779 - learning_rate: 1.0000e-06

Epoch 4/10

66/66 ————— 0s 336ms/step - accuracy: 0.9170 - loss: 0.2163

Epoch 4: val_loss improved from 0.17786 to 0.17734, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 27s 403ms/step - accuracy: 0.9169 - loss: 0.2163 - val_accuracy: 0.9400 - val_loss: 0.1773 - learning_rate: 1.0000e-06

Epoch 5/10

66/66 ————— 0s 339ms/step - accuracy: 0.9253 - loss: 0.1899

Epoch 5: val_loss improved from 0.17734 to 0.17638, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 27s 408ms/step - accuracy: 0.9252 - loss: 0.1901 - val_accuracy: 0.9378 - val_loss: 0.1764 - learning_rate: 1.0000e-06

Epoch 6/10

66/66 ————— 0s 337ms/step - accuracy: 0.9055 - loss: 0.2120

Epoch 6: val_loss did not improve from 0.17638

66/66 ————— 27s 401ms/step - accuracy: 0.9057 - loss: 0.2119 - val_accuracy: 0.9378 - val_loss: 0.1765 - learning_rate: 1.0000e-06

Epoch 7/10

66/66 ————— 0s 339ms/step - accuracy: 0.9249 - loss: 0.1983

Epoch 7: val_loss improved from 0.17638 to 0.17577, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 27s 407ms/step - accuracy: 0.9250 - loss: 0.1981 - val_accuracy: 0.9378 - val_loss: 0.1758 - learning_rate: 1.0000e-06

Epoch 8/10

66/66 ————— 0s 338ms/step - accuracy: 0.9181 - loss: 0.2082

Epoch 8: val_loss improved from 0.17577 to 0.17527, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 27s 406ms/step - accuracy: 0.9181 - loss: 0.2081 - val_accuracy: 0.9356 - val_loss: 0.1753 - learning_rate: 1.0000e-06

Epoch 9/10

66/66 ————— 0s 338ms/step - accuracy: 0.9169 - loss: 0.1955

Epoch 9: val_loss improved from 0.17527 to 0.17522, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 27s 406ms/step - accuracy: 0.9168 - loss: 0.1957 - val_accuracy: 0.9356 - val_loss: 0.1752 - learning_rate: 1.0000e-06

Epoch 10/10

66/66 ————— 0s 337ms/step - accuracy: 0.9253 - loss: 0.1949

Epoch 10: val_loss improved from 0.17522 to 0.17508, saving model to best_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

66/66 ————— 27s 404ms/step - accuracy: 0.9253 - loss: 0.1948 - val_accuracy: 0.9356 - val_loss: 0.1751 - learning_rate: 1.0000e-06
Restoring model weights from the end of the best epoch: 10.

8. Evaluate and Visualize

- Plot training/validation accuracy and loss.
- Show confusion matrix and classification report.
- Visualize predictions on test images.

```
from sklearn.metrics import confusion_matrix, classification_report
import numpy as np
import seaborn as sns

# Plot training/validation accuracy and loss
fig, axs = plt.subplots(1, 2, figsize=(12, 5))
axs[0].plot(history.history['accuracy'], label='Train Acc')
axs[0].plot(history.history['val_accuracy'], label='Val Acc')
if 'accuracy' in fine_tune_history.history:
    axs[0].plot(np.arange(len(history.history['accuracy']),
len(history.history['accuracy']) +
len(fine_tune_history.history['accuracy'])),
fine_tune_history.history['accuracy'], label='Train Acc (ft)')
    axs[0].plot(np.arange(len(history.history['val_accuracy']),
len(history.history['val_accuracy']) +
len(fine_tune_history.history['val_accuracy'])),
fine_tune_history.history['val_accuracy'], label='Val Acc (ft)')
axs[0].set_title('Accuracy')
axs[0].legend()
axs[1].plot(history.history['loss'], label='Train Loss')
axs[1].plot(history.history['val_loss'], label='Val Loss')
if 'loss' in fine_tune_history.history:
    axs[1].plot(np.arange(len(history.history['loss']),
len(history.history['loss']) +
len(fine_tune_history.history['loss'])),
fine_tune_history.history['loss'], label='Train Loss (ft)')
    axs[1].plot(np.arange(len(history.history['val_loss']),
len(history.history['val_loss']) +
len(fine_tune_history.history['val_loss'])),
fine_tune_history.history['val_loss'], label='Val Loss (ft)')
axs[1].set_title('Loss')
axs[1].legend()
plt.show()
```

```

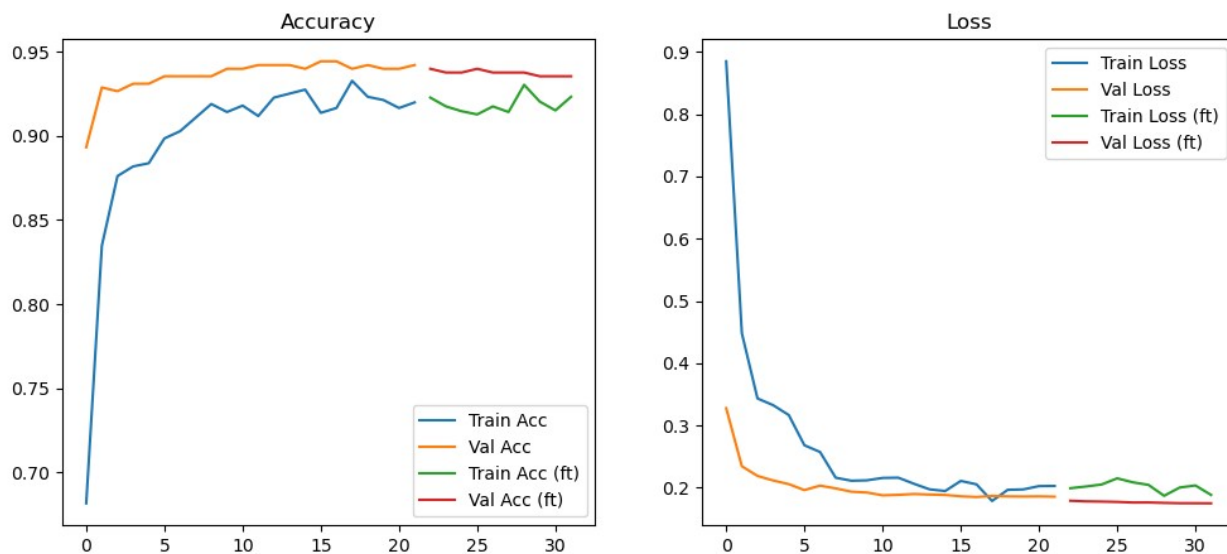
# Predict on test set
y_true = test_gen.classes
y_pred_probs = model.predict(test_gen)
y_pred = np.argmax(y_pred_probs, axis=1)

# Confusion matrix
cm = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=classes, yticklabels=classes)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

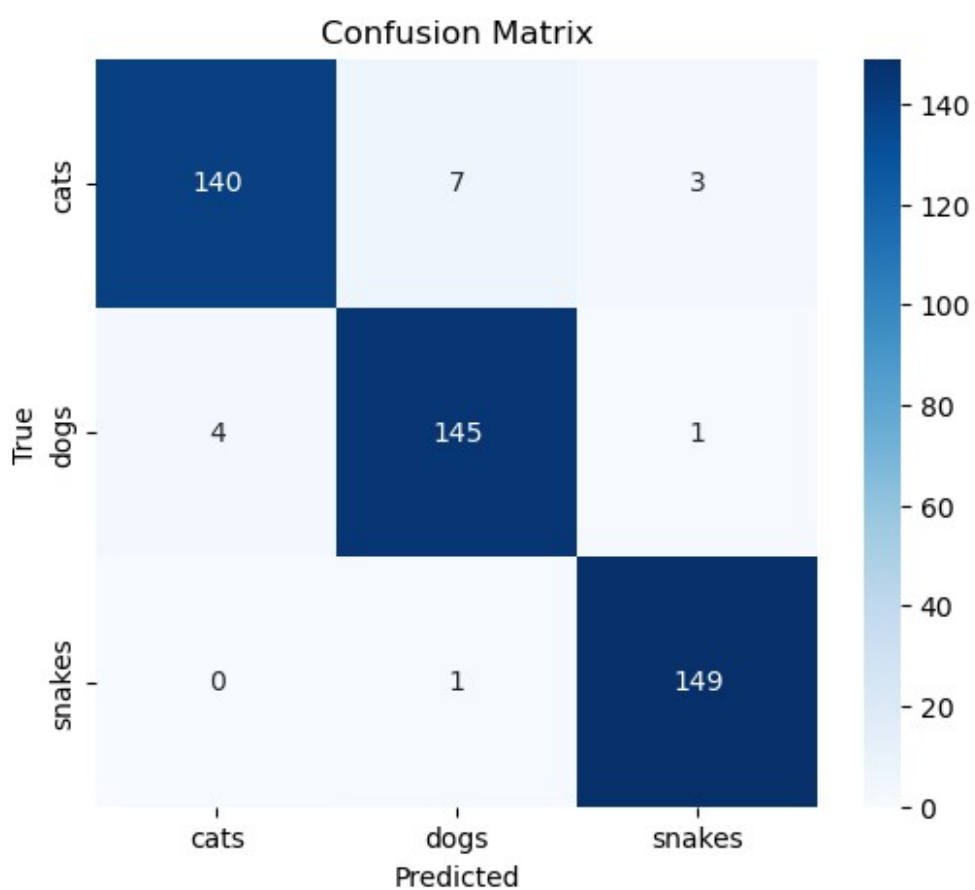
# Classification report
print("Classification Report:")
print(classification_report(y_true, y_pred, target_names=classes))

# Visualize predictions on test images
test_gen.reset()
fig, axes = plt.subplots(2, 5, figsize=(18, 8))
for i, ax in enumerate(axes.flat):
    img, label = next(test_gen)
    pred_idx = np.argmax(model.predict(img), axis=1)[0]
    true_idx = np.argmax(label, axis=1)[0]
    ax.imshow(img[0].astype('uint8'))
    ax.set_title(f"True: {classes[true_idx]}\nPred:
{classes[pred_idx]}")
    ax.axis('off')
plt.tight_layout()
plt.show()

```



15/15 ————— 7s 356ms/step



Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

cats	0.97	0.93	0.95	150
dogs	0.95	0.97	0.96	150
snakes	0.97	0.99	0.98	150
accuracy			0.96	450
macro avg	0.96	0.96	0.96	450
weighted avg	0.96	0.96	0.96	450

```

1/1 ===== 2s 2s/step
1/1 ===== 0s 284ms/step
1/1 ===== 0s 291ms/step
1/1 ===== 0s 319ms/step
1/1 ===== 0s 289ms/step
1/1 ===== 0s 299ms/step
1/1 ===== 0s 285ms/step
1/1 ===== 0s 283ms/step
1/1 ===== 0s 280ms/step
1/1 ===== 0s 282ms/step

```



9. Save and Export the Model

```
# Save the trained model in HDF5 format
model.save('final_model.h5')
```

```
print("Model saved as 'final_model.h5'")
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

Model saved as 'final_model.h5'